

SysV AMD64 ABI Example

```
int auth( const char * user) {  
    size_t i;  
    char buf[16];  
    strncpy(buf, user, sizeof (buf));
```

auth:

push rbp	; save previous frame pointer
mov rbp, rsp	; set new frame pointer
sub rsp, 0x30	; allocate space for locals (i, buf)
movabs rdx, 0x10	; move sizeof(buf) to rdx
lea rax, [rbp-0x20]	; get the address of buf on the stack
mov qword [rbp-0x08], rdi	; move user pointer into stack
mov rsi, qword [rbp-0x08]	; move user pointer back into rsi
mov rdi, rax	; move buf into rdi
call strncpy	; call strncpy(rdi, rsi, rdx)
...	

Writing in Assembly

```
> yasm -- version
```

```
yasm 1.3.0
```

- Ok, enough review, let's write a program
- In keeping with the slides, we'll use an Intel syntax assembler called yasm
 - nasm is equivalent for this course
 - Feel free to use gas if you can't stand Intel syntax
- Let's write the simplest possible program
 - Immediately exit with status code 0

Hello World in Assembly

```
bits 64          ; we are writing a 64-bit program
section .text    ; we will place this in the .text (code) section

extern _exit     ; we are referencing an external function (exit)
                 ; libc functions are prefixed with '_'

global _start    ; declare _start as global symbol
                 ; this preserves a symbol table entry
_start:         ; _start is the default ELF entry point
    mov rdi, 0x00 ; zero out rdi, our first argument
    call _exit    ; call exit(rdi=0)
    int3         ; raise a breakpoint trap if we get here
                 ; (we should never get here)
```

Assembling and Linking

```
> yasm -f elf64 -o exit.o exit.asm
> ld -o exit exit.o -lc
> /lib/ld-2.18.so ./exit
> echo $?
0
```

1. We first assemble the program to an object file exit.o
2. We link an ELF exe against libc
3. We run it using a given runtime loader
 - You might need to specify a different path
 - Or, you might not need to specify it on your system
4. It returns 0!

Disassembly

Disassembling is the process of recovering assembly from machine code

- Not to be confused with decompilation!
- Requires knowledge of binary format and ISA

Distinction between linear sweep and recursive descent disassembly

- Linear sweep begins at an address and continues sequentially until the buffer is exhausted
- Recursive descent disassembly begins at an address and follows program control flow, discovering all reachable code

Disassembly with objdump

```
> objdump -d -M intel exit
exit:      file format elf64-x86-64
```

...

Disassembly of section .text:

```
000000000400230 <_start>:
  400230:  48 c7 c7 00 00 00 00  mov    rdi,0x0
  400237:  e8 e4 ff ff ff      call   400220 <_exit@plt>
  40023c:  cc                  int3
```

Dumping ELF Objects

```
> objdump -x exit
```

```
exit:      file format elf64-x86-64
```

```
[ ... ]
```

```
architecture: i386:x86-64, flags 0x00000112:
```

```
EXEC_P, HAS_SYMS, D_PAGED
```

```
start address 0x0000000000400230
```

```
Program Header:
```

```
[ ... ]
```

```
Dynamic Section:
```

```
    NEEDED                libc.so.6
```

```
[ ... ]
```

```
Sections:
```

```
Idx Name Size VMA LMA File off  Algn
```

```
[ ... ]
```

```
SYMBOL TABLE:
```

```
[ ... ]
```

```
0000000000000000      F *UND* 0000000000000000      _exit@@GLIBC_2.2.5
```

```
[ ... ]
```

Invoking Syscalls

We can, of course, bypass libc and directly ask the kernel for services

- Need to know the syscall number (index into syscall table in the kernel)

On Linux/x86_64, we use the *syscall* instruction to transfer control to the kernel

- On linux/x86, traditionally INT 0x80 (interrupt 128)

Let's write a program to print a message and exit cleanly

Invoking Syscalls

```
bits 64                                ; as before
section .text

global _start

_start:
    mov rdx, msg_len                    ; len(msg) to rdx
    mov rsi, msg                        ; msg to rsi
    mov rdi, 1                          ; fd 1 (stdout) to rdi
    mov rax, 1                          ; write is syscall 1
    syscall                             ; call write(rdi, rsi, rdx)

    mov rdi, 0                          ; status code 0 to rdi
    mov rax, 60                         ; exit is syscall 60
    syscall                             ; call exit(rdi)

section .data                          ; the program's .data section
msg:      db 'aha',0x0a                 ; the message as a byte array
msg_len:  equ $-msg                     ; len is current addr - msg
```

Invoking Syscalls

```
> yasm -f elf64 -o hello.o hello.asm
> ld -o hello hello.o
> ./hello
aha
> echo $?
0
```

- You should see something similar to the above
 - Notice we didn't link against libc – no need
 - And, we do not need to invoke the runtime loader

Debugging Programs

Sometimes your program (or exploit) is buggy

- Well, mine are

An interactive debugger is an invaluable tool in these cases

- Start, stop execution
- Set breakpoints, watchpoints
- Directly inspect memory and CPU state
- Modify program state

Debugging binary programs? gdb!

GDB Quick Reference Card

GDB QUICK REFERENCE GDB Version 4

Essential Commands

`gdb program [core]` debug *program* [using `coredump core`]
`b [file:]function` set breakpoint at *function* [in *file*]
`run [arglist]` start your program [with *arglist*]
`bt` backtrace: display program stack
`p expr` display the value of an expression
`c` continue running your program
`n` next line, stepping over function calls
`s` next line, stepping into function calls

Starting GDB

`gdb` start GDB, with no debugging files
`gdb program` begin debugging *program*
`gdb program core` debug `coredump core` produced by *program*
`gdb --help` describe command line options

Stopping GDB

`quit` exit GDB; also `q` or EOF (eg `C-d`)
`INTERRUPT` (eg `C-c`) terminate current command or

Breakpoints

`break [file:]line`
`b [file:]line`
`break [file:]line`
`break +c`
`break -c`
`break *c`
`break`
`break ..`
`cond n [condition]`

`tbreak .`
`rbreak r`
`watch expr`
`catch expr`

`info breakpoints`
`info watchpoints`

`clear`
`clear [file:]line`
`clear [file:]line`

Get it! from the web!

Initializing GDB

```
> cat .gdbinit  
set disassembly-flavor intel  
disp/i $rip
```

Initial setup

- Set default disassembly syntax
- Display current instruction at each prompt

Also useful for

- Setting breakpoints
- Scripting execution of program to known bad state
- etc.

Let's debug our hello world program from before

Starting GDB

```
> gdb hello
```

```
GNU gdb (Debian 7.7.1+dfsg-3) 7.7.1
```

```
[ ... ]
```

```
(gdb) b _start
```

```
Breakpoint 1 at 0x4000b0
```

- **We load the program in gdb and set a breakpoint at the entry point (`_start`)**
- **Breakpoints insert (by default) a software interrupt (`int3`) at the given address**
- **When `int3` executes, control transfers to gdb**
 - Replaces the original instruction
 - After original instruction is executed, the `int3` is restored

Running the Program

```
(gdb) r
```

```
Starting program: [...]hello
```

```
Breakpoint 1, 0x00000000004000b0 in _start ()
```

```
2: x/i $rip
```

```
=> 0x4000b0 <_start>:    movabs rdx,0x4
```

```
(gdb)
```

- We run the program and immediately hit our breakpoint
- Our display command prints the current instruction

Single-Stepping

```
(gdb) si
0x00000000004000ba in _start ()
2: x/i $rip
=> 0x4000ba <_start+10>:    mov     rsi,0x6000e4
(gdb)
0x00000000004000c1 in _start ()
2: x/i $rip
=> 0x4000c1 <_start+17>:    mov     rdi,0x1
(gdb)
0x00000000004000c8 in _start ()
2: x/i $rip
=> 0x4000c8 <_start+24>:    mov     rax,0x1
(gdb)
0x00000000004000cf in _start ()
2: x/i $rip
=> 0x4000cf <_start+31>:    syscall
```

- We can single-step the program by issuing `si`

Inspecting State

```
(gdb) p $rax
$1 = 1
(gdb) p $rsi
$2 = 6291684
(gdb) p/x $rsi
$3 = 0x6000e4
(gdb) x/s $rsi
0x6000e4:  "aha\n"
```

- p (print) prints register contents
- x (examine) dereferences addresses
- Formatting suffixes control how values are interpreted
 - /x for hex
 - /s for null-terminated strings

More Inspection

```
(gdb) x/8xb $rsp
```

```
0x7fffffffefe790: 0x01      0x00      0x00      0x00      0x00      0x00  
0x00      0x00
```

```
(gdb) x/8xg $rsp
```

```
0x7fffffffefe790: 0x0000000000000001  0x00007fffffffefa5f  
0x7fffffffefe7a0: 0x0000000000000000  0x00007fffffffefa9f  
0x7fffffffefe7b0: 0x00007fffffffefaaa  0x00007fffffffeface  
0x7fffffffefe7c0: 0x00007fffffffefadf  0x00007fffffffefaf2
```

- We can print the stack by referencing \$rsp
- Repetition suffix (above, 8)
- Additional width suffix
 - b for bytes
 - w for dwords
 - g for qwords

Dumping CPU State

```
(gdb) info registers
rax          0x1  1
rbx          0x0  0
rcx          0x0  0
rdx          0x4  4
rsi          0x6000e4 6291684
rdi          0x1  1
rbp          0x0  0x0
rsp          0x7fffffffef790 0x7fffffffef790
r8           0x0  0
[ ... ]
r15          0x0  0
rip          0x4000cf 0x4000cf <_start+31>
eflags       0x202  [ IF ]
cs           0x33  51
ss           0x2b  43
[ ... ]
```

- We can dump the entire (visible) CPU state

Mutating State

```
(gdb) set $r10=0x31337
```

```
(gdb) p/x $r10
```

```
$1 = 0x31337
```

```
(gdb) x/s $rsi
```

```
0x6000e4: "aha\n"
```

```
(gdb) set *0x6000e4=0x0a485542
```

```
(gdb) x/s $rsi
```

```
0x6000e4: "BUH\n"
```

- We can set values both in registers and in memory

Continuing Execution

```
(gdb) c
```

```
Continuing.
```

```
BUH
```

```
[Inferior 1 (process 7228) exited normally]
```

- And we can continue execution
 - Here, we exit since we don't hit our breakpoint again
 - And, we print out our modified string

Summary

You should now be able to

1. Follow and understand simple assembly programs
2. Create your own programs
3. Debug them using gdb

We are only scratching the surface, but it's enough to get started